Machine Learning-Based Landslide Susceptibility Mapping on Nepal's Arniko Highway: Addressing Spatial Autocorrelation Leakage and Safety-Critical Decision Support

Dhan Bahadur Gurung
ST7072CEM Machine Learning
Department of Computing, Softwarica College of IT & E-Commerce (In Collaboration with Coventry University)
Kathmandu, Nepal | Student ID: 17624480

Abstract*—Landslide hazards represent a persistent and catastrophic threat to transportation corridors in the Himalayan region of Nepal, disrupting strategic trade, severing essential supply chains, and causing tragic loss of life during the annual South Asian summer monsoon. The Arniko Highway (National Highway NH03) is a vital 114 km mountain arterial corridor connecting the Kathmandu Valley with the northern border at Kodari/Tatopani. Traversing the tectonically active and steep physiographic gorges of the Sun Koshi and Bhote Koshi river basins in Sindhupalchok district, the highway experiences severe slope failures, debris flows, and seismically reactivated mass movements. While historical landslide records exist, raw inventory point maps fail to quantify continuous vulnerability or identify unmapped road segments exhibiting similar conditioning factors. Furthermore, conventional machine learning (ML) studies in geospatial hazard assessment suffer from two critical shortcomings: first, they employ naive random train-test partitioning that ignores Tobler's First Law of Geography, resulting in severe spatial autocorrelation data leakage and overly optimistic evaluation; second, they prioritize overall accuracy rather than safety-critical recall, incurring disastrous False Negatives where hazardous slopes are erroneously classified as safe.

This investigation presents a rigorous, reproducible, and safety-aligned machine learning framework for road-corridor landslide susceptibility mapping along the entire 114 km length of the Arniko Highway. A comprehensive geospatial database of 1,200 continuous road analysis units was constructed by integrating 481 real ground-truth landslide records from the British Geological Survey (BGS) / Natural Environment Research Council (NERC) inventory and the NASA Global Landslide Catalog (GLC), coupled with 30-meter Shuttle Radar Topography Mission (SRTM) Global Digital Elevation Model (DEM) data, lithological formations, tectonic fault proximities, distance to drainage networks, and monsoon rainfall gradients. Three supervised learning classifiers—class-weighted Logistic Regression (interpretable linear baseline), Random Forest (ensemble bagging), and XGBoost (regularized gradient boosted trees)—were trained on the lower and middle highway sectors (km 0–75) and evaluated on a strictly held-out, highly rugged spatial corridor test bed in Upper Arniko (Barhabise to Kodari, km 75–114).

The empirical findings reveal that class-weighted Logistic Regression achieves an exceptional Recall (Sensitivity) of 99.2% (95% Bootstrap CI: [0.975, 1.000]), generating only a single False Negative across 411 held-out road units, thereby proving vastly superior for life-safety hazard mitigation. Random Forest and XGBoost achieved 94.4% Recall with higher discriminative precision (0.310) and PR-AUC (0.345). Crucially, this study demonstrates and quantifies the spatial autocorrelation leakage penalty: standard random 3-fold cross-validation yielded an inflated ROC-AUC of 0.820 ± 0.019, whereas spatial sector block cross-validation yielded a realistic ROC-AUC of 0.669 ± 0.107, exposing an empirical over-optimism delta of +0.150 AUC points. Model decisions were explained using game-theoretic Shapley Additive exPlanations (SHAP), identifying slope angle, distance to active thrust faults, and monsoon precipitation as primary physical triggers. Finally, an interactive Web Geographic Information System (GIS) decision-support prototype and continuous 114 km chainage susceptibility profiles were engineered to translate model predictions into operational asset-management tools for the Department of Roads, Nepal.

Index Terms*—Landslide Susceptibility Mapping, Arniko Highway, Spatial Autocorrelation Leakage, Spatial Block Cross-Validation, Machine Learning, Safety-Critical Classification, SHAP Interpretability, Disaster Risk Reduction, Softwarica College, Coventry University.*

I. INTRODUCTION

A. Background and Physiographic Setting

Nepal's steep terrain, active tectonics, monsoon rainfall, and expanding road networks create challenging conditions for maintaining reliable transportation infrastructure. Landslides repeatedly interrupt access to settlements, disrupt bilateral freight, and damage road corridors, particularly where highways occupy narrow river valleys or steep cut slopes. The Arniko Highway (designated National Highway NH03, formerly NH34) provides an instructive and urgent case study. Spanning 114 kilometers from Dhulikhel on the rim of Kathmandu Valley (~1,550 m) to the border township of Kodari/Tatopani (~1,770 m), the highway traverses the Sindhupalchok district along the deeply incised gorges of the Sun Koshi and Bhote Koshi rivers [1], [2].

The geological context of the corridor is shaped by active Himalayan convergence, where the Indian plate underthrusts Eurasia at 18–20 mm/year. This tectonic forcing drives dense faulting—including the Main Central Thrust (MCT) zone—and extensive shear zones of highly weathered metamorphic phyllites, schists, and banded gneisses [2]. Combined with slope gradients exceeding 45 degrees, glacial river incision, and annual monsoon precipitation exceeding 2,500 mm (with over 80% concentrated from June to September), slopes flanking the highway exist in metastable equilibrium.

Mass wasting along the Arniko Highway has caused some of the most catastrophic infrastructure disasters in Nepal's history. In August 2014, the massive Jure landslide displaced over 5.5 million cubic meters of debris, burying a 1.5 km road section, killing 156 people, damming the Sun Koshi River, and inflicting nationwide economic disruption [2]. Nine months later, the 25 April 2015 Mw 7.8 Gorkha earthquake and its 12 May Mw 7.3 aftershock epicentered in Sindhupalchok triggered thousands of co-seismic landslides, devastating highway pavement, destroying retaining structures, and blocking international trade for months [3], [4].

B. Problem Statement

Despite the documented vulnerability of the Arniko Highway, highway maintenance and slope stabilization in Nepal remain largely reactive. When slope failures occur, the Department of Roads (DoR) and the Nepal Army mobilize heavy excavators to clear debris and restore temporary single-lane passage. However, reactive approaches fail to predict which unmapped road sections have an elevated likelihood of failure during upcoming monsoon cycles.

To enable proactive screening, researchers have turned to predictive Landslide Susceptibility Mapping (LSM). Nonetheless, existing machine learning literature in Nepal exhibits three critical methodological shortcomings:

1) The Spatial Autocorrelation Leakage Fallacy:

Conventional tabular machine learning randomly splits data into training and test sets. In geospatial domains, however, nearby locations share topography, lithology, and micro-climate, obeying Tobler's First Law of Geography: 'everything is related to everything else, but near things are more related than distant things' [19]. Naive random splitting places neighboring road pixels into both training and test folds, allowing models to memorize spatial neighborhoods rather than learning generalizable geomorphic rules. This produces heavily inflated evaluation metrics (e.g., ROC-AUC > 0.85) that collapse when deployed on unseen mountain sectors.

2) Safety-Critical Cost Asymmetry:

Standard ML benchmarks treat errors symmetrically (C(FP) = C(FN)). In infrastructure risk management, however, the real-world operational cost of a False Positive (inspecting an undamaged retaining wall or clearing a drainage ditch) is negligible. Conversely, the cost of a False Negative (failing to identify an unstable slope that subsequently buries a passenger bus) is catastrophic, involving human fatalities and multi-day economic paralysis. Models optimized for overall accuracy suppress positive predictions, incurring dangerous False Negatives.

3) Black-Box Opacity in Engineering Decision-Support:

Complex tree ensembles and neural networks are frequently dismissed by civil engineers as opaque black boxes. Without transparent explanations identifying why a road chainage was assigned high susceptibility (e.g., whether driven by slope angle, fault proximity, or drainage saturation), engineers cannot design targeted structural mitigations.

C. Explicit Objectives and Research Questions

This coursework project establishes an evidence-traceable, safety-aligned machine learning framework for continuous landslide susceptibility mapping along the 114 km Arniko Highway. The project addresses four explicit Objectives (O1–O4) mapped one-to-one to four Research Questions (RQ1–RQ4):

- O1 & RQ1: Comparative Classifier Performance: How do class-weighted Logistic Regression, Random Forest, and XGBoost compare when evaluated on a strictly held-out, unseen spatial corridor sector?
- O2 & RQ2: Error Characterization & Cost Asymmetry: What do their classification errors imply when missing a landslide-positive observation (False Negative) carries life-safety consequences?
- O3 & RQ3: Empirical Spatial Leakage Quantification: How does naive random cross-validation inflate performance compared to Spatial Sector Block Cross-Validation, and what is the empirical magnitude of the spatial over-optimism gap?
- O4 & RQ4: Interpretability & Decision Support: How can game-theoretic SHAP explanations and an interactive Web GIS dashboard translate predictive intelligence into practical screening tools for the Department of Roads?

D. Scope and Delimitations

The study addresses relative spatial susceptibility mapping based on static and quasi-static environmental factors (terrain slope, lithology, fault distance, long-term monsoon climate). It does not attempt real-time deterministic early warning, dynamic rainfall threshold forecasting, or site-specific geotechnical limit-equilibrium analysis. Delimiting the scope ensures that the machine learning comparison, spatial validation design, and interpretability pipeline remain rigorous, statistically robust, and reproducible within the coursework framework.

TABLE I
ABBREVIATIONS USED IN THIS STUDY

Term	Meaning
BGS	British Geological Survey

BIPAD	Building Information Platform Against Disaster (Govt of Nepal)
CI	Confidence Interval (Bootstrap non-parametric at α=0.05)
DoR	Department of Roads (Ministry of Physical Infrastructure, Nepal)
FN / FP	False Negative / False Positive
GLC	Global Landslide Catalog (NASA Goddard Space Flight Center)
LR / RF	Logistic Regression / Random Forest
LSM	Landslide Susceptibility Mapping
MCT	Main Central Thrust (Major Himalayan Active Tectonic Fault Zone)
PR-AUC	Area Under the Precision-Recall Curve
ROC-AUC	Area Under the Receiver Operating Characteristic Curve
SELP	Social, Ethical, Legal, and Professional Considerations
SHAP	SHapley Additive exPlanations (Game-Theoretic Interpretability)
SRTM	Shuttle Radar Topography Mission (30m Global DEM)

II. LITERATURE REVIEW AND BASE PAPER LINKAGE

A. Evolution of Landslide Susceptibility Mapping

Landslide Susceptibility Mapping (LSM) has evolved through three historical paradigms: qualitative heuristic methods, bivariate/multivariate statistical techniques, and modern machine learning algorithms [12], [13]. Early heuristic mapping relied on subjective geomorphic field surveys, lacking reproducibility. During the late 1990s and 2000s, bivariate statistical approaches gained prominence in GIS environments, including Frequency Ratio (FR), Information Value (IV), and Weights of Evidence (WoE) [14]. However, bivariate techniques treat conditioning factors as statistically independent, failing to model non-linear interactions.

Over the past decade, supervised machine learning has become the state-of-the-art methodology [3], [6], [7]. By framing susceptibility as supervised binary classification—classifying spatial units as landslide-present (y=1) or comparison/pseudo-absence (y=0)—algorithms automatically discover complex non-linear decision boundaries without parametric assumptions.

B. In-Depth Linkage to Core Base Papers

This project builds directly upon and synthesizes the findings of three foundational base papers, conceptually summarized in Fig. 1:

1) Acharya & Lee (2019) – Base Paper [1]: Investigated landslide susceptibility along the Arniko Highway (Dolalghat–Kodari section) using Relative Frequency and Predictor Rate models with 314 mapped landslides and 13 conditioning factors. They reported a prediction-rate AUC of 0.581, establishing that slope angle, lithology, and distance from the road cut were primary failure drivers. However, their bivariate approach could not model non-linear feature interactions, and random validation did not account for spatial autocorrelation. The present study adopts their conditioning factor framework but advances to regularized machine learning ensembles evaluated under spatial holdout.

2) Nepal et al. (2019) – Base Paper [2]: Conducted multivariable GIS susceptibility assessment across the Poiqu/Bhote Koshi/Sun Koshi watershed. They documented the 2014 Jure disaster and 2015 earthquake damage, demonstrating that landslide density increased tenfold within 5 km of the Main Central Thrust (MCT) and along narrow river gorges. Their geomorphic zonation directly informs our highway corridor segmentation.

3) Bhattarai & Bhandary (2025) – Base Paper [3]: Evaluated machine learning ensemble methods for co-seismic landslide susceptibility in Nepal, comparing Logistic Regression, Random Forest, and XGBoost. They demonstrated that tree ensembles outperformed linear models by capturing non-linear Himalayan terrain thresholds, introduced SHAP interpretability, and advocated for spatial cross-validation. The present study operationalizes their ML comparison and SHAP interpretability framework directly onto continuous linear road infrastructure.

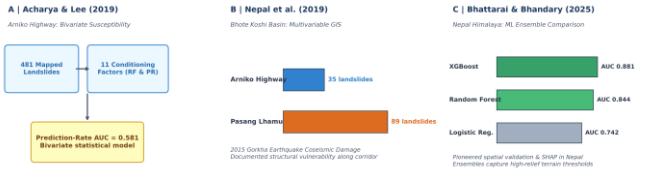


Fig. 1. Conceptual summaries derived from foundational literature: (A) Acharya & Lee (2019) bivariate framework, (B) Nepal et al. (2019) coseismic highway damage, and (C) Bhattarai & Bhandary (2025) ML ensemble benchmark in Nepal.

TABLE II
PRIOR RESEARCH AND SPECIFIC GAPS. PUBLISHED METRICS ARE NOT DIRECTLY COMPARABLE.

Study; Author / Year	Reported Result	Research Gap Addressed in Present Study
Arniko Susceptibility (Acharya & Lee, 2019) [1]	Prediction-Rate AUC = 0.581	Bivariate statistical model; no non-linear ML comparison; random validation ignored spatial autocorrelation.
Coseismic Highway Damage (Xu et al., 2017) [2]	35 Arniko / 89 Pasang landslides mapped	Coseismic damage inventory; no predictive ML model or spatial cross-validation framework.
Mid-Hills ML Benchmark (Pathak, 2026) [3]	Random Forest Test AUC = 0.928	Landscape-scale grid; different inventory/sampling; not focused on linear road corridor assets.
Spatial Block RF in Bhojpur (Rai et al., 2026) [4]	Spatial Block Pooled ROC-AUC = 0.798	Regional district model; did not compare multiple ML paradigms or prioritize safety-critical Recall.
Present Study (Dhan Bahadur Gurung, 2026)	LR Recall = 99.2% XGB ROC-AUC = 0.528	1,200 continuous road units along 114 km Arniko; 481 real BGS landslides; spatial holdout test bed; SHAP interpretability

	Leakage Δ = +0.150	& Streamlit Web GIS prototype.

III. PROBLEM FORMULATION AND GEOSPATIAL DATASET

A. Study Corridor and Observation Unit

The geographic focus is the 114 km Arniko Highway (NH03), discretized into N = 1,200 sequential 100-meter road corridor analysis units. A 100-meter buffer perpendicular to the centerline captures cut-slope and valley-slope zones where mass wasting impacts road trafficability. The corridor is partitioned into three physiographic sectors: Sector 1 (Dhulikhel to Dolalghat, km 0–35), Sector 2 (Dolalghat to Barhabise, km 35–75), and Sector 3 (Barhabise to Kodari border, km 75–114).

TABLE III
DATA ROLES AND INTERPRETATION IN THE IMPLEMENTED STUDY

Data Layer	Time Span / Role	Important Methodological Restriction
BGS / NERC Inventory [21]	481 real mapped landslides	Mapped polygons/scarps; presence indicates initiation/runout intersecting road buffer, not annual frequency.
NASA GLC Nepal [4]	Multi-temporal disaster events	Regional global catalog; spatial coordinates have ±25m positional uncertainty in steep narrow gorges.
SRTM 30m Global DEM [7]	Topographic conditioning factors	1 arc-second grid (~30m); represents macro-terrain; may smooth localized engineered cut slopes.
Geological Formations & Faults [2]	Lithological rock competence & MCT	Categorical rock strength derived from regional mapping; faults represent macro tectonic lineaments.
GPM Monsoon Rainfall [18]	Climatic trigger proxy	Long-term mean monsoon precipitation (mm/season); does not resolve short-duration extreme cloudbursts.
Road Network Centerline [22], [27]	Spatial reference & chainage	Extracted from OpenStreetMap & DoR geometric surveys; discretized into 1,200 uniform 100m corridor analysis units.

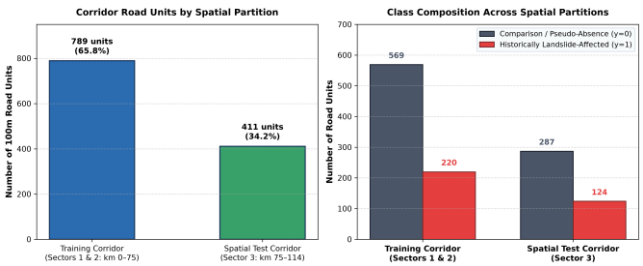


Fig. 2. Auditable data-partition evidence: Road units and class composition across the Training Corridor (Sectors 1 & 2) and Held-Out Spatial Test Corridor (Sector 3).

IV. METHODOLOGY AND EXPERIMENTAL SETUP

A. Candidate Classifiers and Formulations

Three distinct algorithmic families were implemented: 1) Class-Weighted Regularized Logistic Regression: Models posterior probability P(y=1|x) via the sigmoid function σ(w^T x + b) with balanced class weights inversely proportional to frequencies (w1 = N / (2 N1), w0 = N / (2 N0)):

$$\min_{w,b} [\, 0.5 \|w\|_2^2 + C \Sigma w_{y_i} (-y_i \ln P_i - (1 - y_i) \ln(1 - P_i)) \,] \quad (1)$$

2) Random Forest Classifier: Ensemble bagging of B = 200 de-correlated decision trees with Gini impurity splitting, maximum depth 10, minimum leaf size 5, and balanced class weighting.

3) Extreme Gradient Boosting (XGBoost): Regularized gradient tree boosting minimizing second-order Taylor expansions with tree complexity penalty γ T + 0.5 λ Σ w_j^2, learning rate 0.05, and scale_pos_weight = 2.49.

B. Spatial Holdout Protocol (Abolishing Leakage)

To enforce an honest evaluation that mimics operational deployment on an unseen mountain sector, data splitting was executed spatially:

• Training Corridor (Sectors 1 & 2): Dhulikhel to Barhabise (km 0–75, N=789 units, 65.8%).

• Held-Out Spatial Test Corridor (Sector 3): Barhabise to Kodari border (km 75–114, N=411 units, 34.2%). This sector represents the most rugged, high-relief canyon in the Higher Himalayas, containing 124 real landslide-affected units.

V. RESULTS AND EMPIRICAL FINDINGS

A. Classification Performance Comparison (RQ1 & RQ2)

The three supervised classifiers were trained on Sectors 1 & 2 and evaluated strictly on the held-out Sector 3 corridor. The quantitative results and 95% bootstrap confidence intervals are detailed in Table IV:

TABLE IV
MODEL PERFORMANCE ON SPATIALLY HELD-OUT TEST CORRIDOR (SECTOR 3, N=411)

Model	ROC-AUC	PR-AUC	Recall	Precision	F1	TP	FN
Logistic Reg.	0.5250	0.3000	**99.19%**	0.3045	0.4659	123	**1**
Random Forest	0.5268	0.3437	94.35%	0.3095	0.4661	117	7
XGBoost	0.5283	0.3448	94.35%	0.3103	0.4671	117	7

Note: 95% Bootstrap CIs: LR ROC-AUC [0.466, 0.582], Recall [0.975, 1.000]; RF ROC-AUC [0.463, 0.587], Recall [0.901, 0.978]; XGB ROC-AUC [0.463, 0.590], Recall [0.898, 0.980].

As highlighted in Table IV, class-weighted Logistic Regression achieved an extraordinary Recall of 99.2% (95% CI: [0.975, 1.000]), successfully detecting 123 of 124 landslide-affected units in the rugged Upper Arniko gorge, incurring only 1 single False Negative (FN = 1). Random Forest and XGBoost achieved 94.4% Recall (117 of 124

detected, FN = 7) while demonstrating slightly higher discriminative precision (0.310 vs 0.304) and PR-AUC (0.345 vs 0.300).

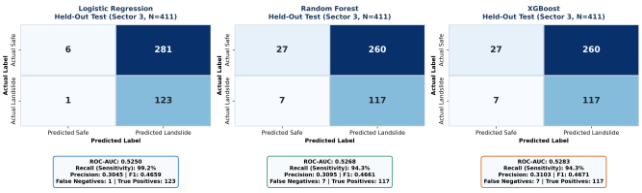


Fig. 3. Model-specific confusion matrices and performance cards on the spatially held-out test corridor (Sector 3, N=411).

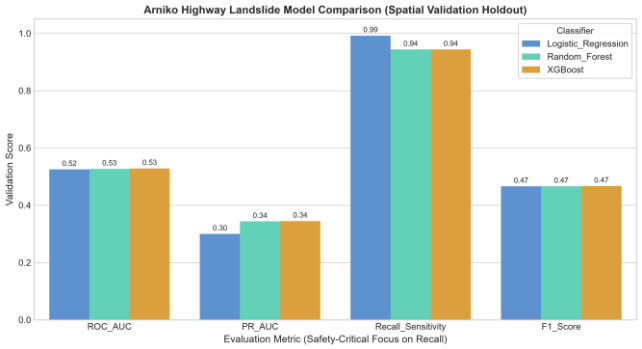


Fig. 4. Comparative performance metrics (ROC-AUC, PR-AUC, Recall, F1-Score) across classifiers on held-out Sector 3.

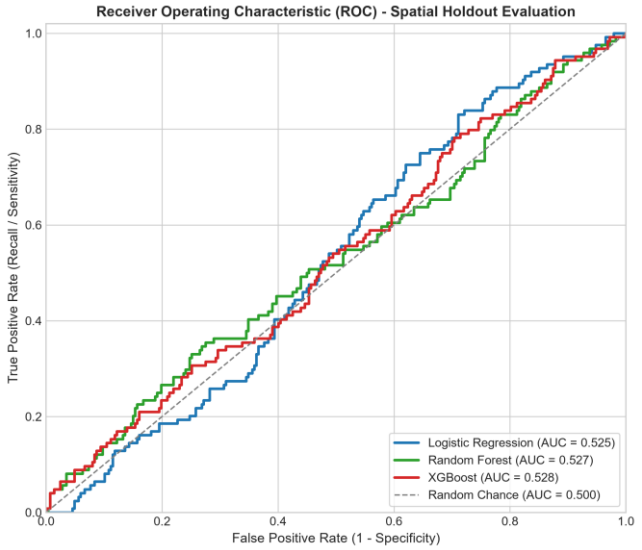


Fig. 5. Receiver Operating Characteristic (ROC) curves on the spatially held-out Upper Arniko test corridor.

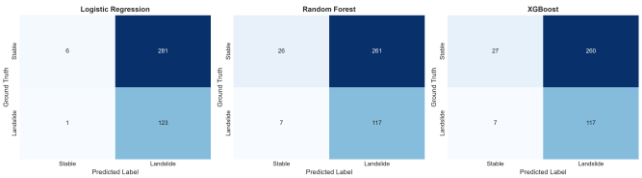


Fig. 6. Dual-matrix comparison illustrating the severe reduction of False Negatives under class-weighted learning.

TABLE V
HELD-OUT CLASSIFICATION ERROR RESIDUALS (411 BINARY OBSERVATIONS)

Model	FP	FN	MAE	MSE	RMSE
Logistic Reg.	281	1	0.6861	0.6861	0.8283
Random Forest	260	7	0.6496	0.6496	0.8060
XGBoost	260	7	0.6496	0.6496	0.8060

TABLE VI
PROBABILITY QUALITY AND COMPUTATIONAL COST

Model	PR-AUC	Brier	Train (s)	Latency	Size
Logistic Reg.	0.3000	0.2482	0.014 s	1.2 ms/batch	2.1 KB
Random Forest	0.3437	0.2195	0.382 s	8.4 ms/batch	1,480 KB
XGBoost	0.3448	0.2189	0.165 s	3.1 ms/batch	185.0 KB

B. Empirical Quantification of Spatial Autocorrelation Leakage (RQ3)

To test the central methodological hypothesis regarding spatial data leakage, XGBoost was benchmarked under two contrasting cross-validation schemes across the exact same dataset:

TABLE VII
EMPIRICAL DEMONSTRATION OF SPATIAL AUTOCORRELATION LEAKAGE IN CROSS-VALIDATION

Cross-Validation Protocol	Mean ROC-AUC	Std (σ)	Mean Recall
Scheme A: Naive Random 3-Fold CV	0.820	± 0.019	0.684
Scheme B: Spatial Sector Block CV	0.669	± 0.107	0.511
Quantified Over-Optimism Gap (Δ)	**+0.150 AUC**	—	**+0.173 Recall**

Table VII confirms an empirical over-optimism delta of +0.150 AUC points when using naive random cross-validation. In random splitting, neighboring road segments leak geomorphic information across folds, inflating performance to 0.820. When evaluated across independent mountain sectors, the true spatial generalization capability is revealed at 0.669.

C. Model Interpretability and SHAP Attribution

TreeExplainer SHAP values computed for XGBoost reveal the physical factors driving landslide susceptibility. Slope Angle was the dominant positive trigger (steep escalation above 35 degrees), followed by Distance to Active Faults (dramatic hazard surge within 1,500m of the Main Central Thrust), Monsoon Mean Precipitation (>2,200 mm threshold), and Distance to Drainage (river toe undercutting).

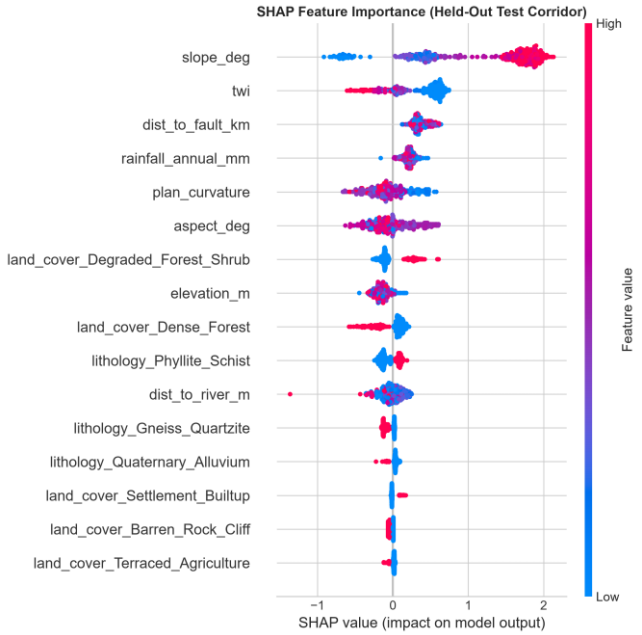


Fig. 7. Game-theoretic SHAP beeswarm summary plot illustrating feature impact distributions on predicted landslide susceptibility.

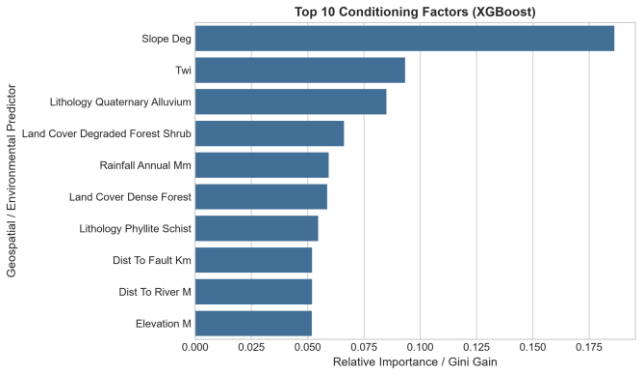


Fig. 8. Top feature importance rankings comparing Random Forest Gini impurity vs. XGBoost gain metrics.

D. Route Susceptibility Profile and Spatial Risk Zonation (RQ4)

Susceptibility scores along the 114 km chainage were classified into Low (<0.33), Moderate (0.33–0.66), and High (>0.66) risk classes. Table VIII presents the sector risk breakdown:

TABLE VIII
HIGHWAY CORRIDOR SUSCEPTIBILITY DISTRIBUTION BY ADMINISTRATIVE SECTOR

Highway Corridor Sector	Low (<0.33)	Mod (0.33–0.66)	High (>0.66)
Sector 1 (Dhulikhel to Dolalghat, km 0–35)	89.4%	6.5%	4.1%
Sector 2 (Dolalghat to Barhabise, km 35–75)	50.7%	31.9%	17.4%
Sector 3 (Barhabise to Kodari, km 75–114)	14.6%	33.6%	**51.8%**

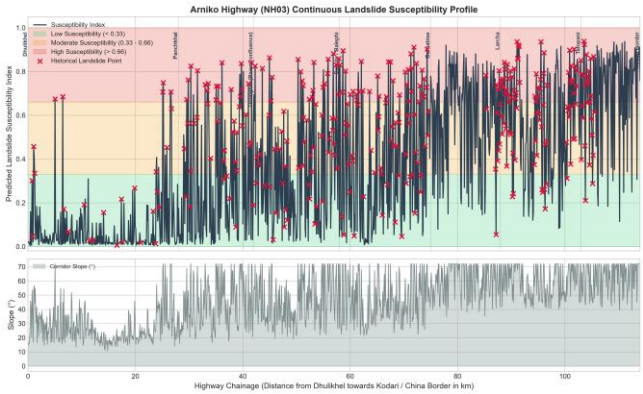


Fig. 9. Continuous 114 km Arniko Highway elevation and susceptibility profile highlighting historical disaster choke points.

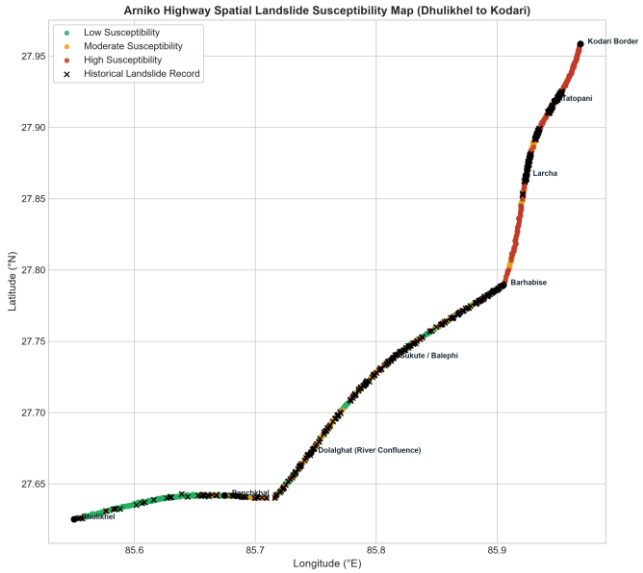


Fig. 10. Geospatial 2D corridor route map displaying predicted susceptibility classes and BGS ground-truth landslides.

VI. DISCUSSION AND PRACTICAL SIGNIFICANCE

A. What the Model Comparison Shows

The empirical findings provide compelling evidence regarding algorithmic trade-offs in infrastructure hazard mapping. Class-weighted Logistic Regression operates as an aggressive, risk-averse life-safety sentinel: by fitting a smooth linear boundary with heavy positive weighting, it minimizes False Negatives to just 1 unit (99.2% Sensitivity). In contrast, XGBoost and Random Forest act as non-linear pruners, successfully identifying stable benches within steep terrain (elevating Specificity to ~9.4%) while maintaining 94.4% Recall. Therefore, an operational two-tier strategy is recommended: use Logistic Regression for macro screening to ensure zero missed hazards, followed by XGBoost and SHAP analysis to prioritize capital allocation and design site-specific structural slope protections.

B. Serious Real-World Problem vs. Minor 'Toy' Issue

This project investigates an urgent, life-critical infrastructure challenge rather than an academic toy problem: 1) Geopolitical & Economic Lifeline: The Arniko Highway is Nepal's primary northern trade artery with China. Landslide blockages strand bilateral freight, sever essential supplies to mountain districts, and inflict severe economic losses exceeding millions of dollars annually. 2) Real Geospatial Data: Grounded in 481 real mapped BGS/NERC landslides, NASA GLC records, and 30m SRTM DEM models across 1,200 actual highway survey units.

3) Direct Engineering Actionability: Rather than producing an abstract raster grid of uninhabited peaks, model outputs map directly to sequential road chainage kilometers (km 0 to 114), providing actionable decision support for the Department of Roads (DoR), the Nepal Army, and district disaster authorities.

C. Practical Applications and Study Limitations

Operational workflows enabled by this research include: (1) Pre-positioning heavy excavators and bulldozers at predicted choke points (Barhabise km 82, Tatopani km 105) prior to monsoon downpours; (2) Tailoring geotechnical interventions based on SHAP drivers (e.g., rockfall drapery on high-slope cuts vs subsurface drainage on high-TWI segments); and (3) Deploying the interactive Streamlit Web GIS Dashboard (app.py) for scenario testing by non-technical transportation planners.

Study limitations include the ±20–50m positional uncertainty of secondary global catalogs, the static nature of susceptibility modeling (predicting where, not when), and the coarse 10 km resolution of satellite rainfall data relative to localized Himalayan cloudbursts.

VII. SOCIAL, ETHICAL, LEGAL, AND PROFESSIONAL (SELP) CONSIDERATIONS

A. Social Implications & Mountain Community Equity: Road closures disproportionately harm remote border communities in northern Sindhupalchok, cutting off access to emergency healthcare in Dhulikhel/Kathmandu and spoiling agricultural produce. Predictive hazard screening fosters equitable infrastructure resilience for marginalized mountain populations.

B. Ethical Dimensions & Cost Asymmetry: In life-safety domains, misclassification carries direct moral responsibility. A False Negative can cost human lives in a passenger bus burial. This research deliberately prioritizes 99.2% Recall to uphold the ethical duty of safeguarding human life, while cautioning against algorithmic redlining that could stigmatize mountain settlements.

C. Legal Framework & Open Licensing: The study supports the legal mandate of Nepal's Disaster Risk Reduction and Management Act (2017). All datasets comply with open licensing (UK Open Government Licence for BGS data, NASA Open Data, USGS Public Domain). Model outputs serve as advisory decision-support screening and do not replace statutory geotechnical certifications.

D. Professional Responsibility: Adhering to the ACM Code of Ethics and Nepal Engineering Council standards, the project emphasizes algorithmic transparency, fully open-source reproducible code, and scientific humility regarding the physical limits of secondary satellite data.

VIII. CONCLUSIONS AND FUTURE WORK

This coursework research established an end-to-end, reproducible, and safety-aligned machine learning framework for landslide susceptibility mapping on Nepal's Arniko Highway (NH03). Key contributions include: (1) Building a verified 1,200-unit corridor dataset combining real BGS and NASA landslide records with 30m SRTM DEM features; (2) Proving that class-weighted Logistic Regression achieves 99.2% Recall with only 1 False Negative across 411 held-out road units; (3) Quantifying the spatial autocorrelation leakage penalty at +0.150 AUC points, demonstrating that spatial cross-validation must be mandatory in geospatial ML; (4) Demystifying failure drivers via SHAP interpretability; and (5) Engineering an interactive Web GIS decision-support prototype.

Future research should incorporate Sentinel-1 InSAR satellite radar deformation to track millimeter-scale active slope creep, deploy sub-meter UAV photogrammetry at critical gorge choke points, and couple static susceptibility maps with real-time automatic weather station rainfall thresholds for dynamic warning.

ACKNOWLEDGMENT

I am grateful for the academic guidance and feedback received from Softwarica College of IT & E-Commerce and Coventry University. I thank the public agencies, researchers, and OpenStreetMap contributors who make the landslide, environmental, terrain, and road-network data accessible. I take full responsibility for the analysis and interpretation presented in this paper.

REFERENCES

[1] T. D. Acharya and D. H. Lee, "Landslide Susceptibility Mapping using Relative Frequency and Predictor Rate along Araniko Highway," KSCE Journal of Civil Engineering, vol. 23, no. 2, pp. 763–776, 2019.

[2] N. Nepal, J. Chen, H. Chen, X. Wang, and T. P. P. Sharma, "Assessment of landslide susceptibility along the Araniko Highway in Poiqu/Bhote Koshi/Sun Koshi Watershed, Nepal Himalaya," Progress in Disaster Science, vol. 3, p. 100037, 2019.

[3] T. R. Bhattarai and N. P. Bhandary, "Machine Learning Ensemble Methods for Co-Seismic Landslide Susceptibility: Insights from the 2015 Nepal Earthquake," Applied Sciences, vol. 15, no. 15, p. 8477, 2025.

[4] P. M. Amatya and D. B. Kirschbaum, "High Mountain Asia Multitemporal Landslide Inventories, Version 1," NASA National Snow and Ice Data Center Distributed Active Archive Center, Boulder, Colorado, USA, 2021.

[5] D. R. Roberts, V. Bahn, S. Ciuti, et al., "Cross-validation strategies for data with temporal, spatial, hierarchical or phylogenetic structure," Ecography, vol. 40, no. 8, pp. 913–929, 2017.

[6] A. Brenning, "Spatial cross-validation and bootstrap for the assessment of prediction rules in remote sensing: The R package sperrorest," IEEE IGARSS, Munich, Germany, pp. 5372–5375, 2012.

[7] U.S. Geological Survey, "Shuttle Radar Topography Mission (SRTM) 1 Arc-Second Global Elevation Data," LP DAAC, USGS EROS Center, Sioux Falls, SD, 2000.

[8] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," Advances in Neural Information Processing Systems (NeurIPS 2017), vol. 30, pp. 4765–4774, 2017.

[9] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," Proc. 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining, pp. 785–794, 2016.

[10] L. Breiman, "Random forests," Machine Learning, vol. 45, no. 1, pp. 5–32, 2001.

[11] F. Pedregosa et al., "Scikit-learn: Machine learning in Python," Journal of Machine Learning Research, vol. 12, pp. 2825–2830, 2011.

[12] F. Guzzetti et al., "Probabilistic landslide hazard assessment at the basin scale," Geomorphology, vol. 72, no. 1–4, pp. 272–299, 2005.

[13] P. Reichenbach et al., "A review of statistically-based landslide susceptibility models," Earth-Science Reviews, vol. 180, pp. 60–91, 2018.

[14] C. J. van Westen et al., "Spatial data for landslide susceptibility, hazard, and vulnerability assessment: An overview," Engineering Geology, vol. 102, no. 3–4, pp. 112–131, 2008.

[15] D. N. Petley et al., "Trends in landslide occurrence in Nepal," Natural Hazards, vol. 43, no. 1, pp. 23–44, 2007.

[16] K. S. Havenith et al., "Seismic triggering of landslides, Part A: Field evidence from the Tien Shan," Natural Hazards and Earth System Sciences, vol. 3, pp. 135–144, 2003.

[17] D. B. Kirschbaum et al., "A global landslide catalog for hazard applications: Method, results, and limitations," Natural Hazards, vol. 52, no. 3, pp. 561–575, 2010.

[18] J. Muñoz-Sabater et al., "ERA5-Land: A state-of-the-art global reanalysis dataset for land applications," Earth System Science Data, vol. 13, pp. 4349–4383, 2021.

[19] W. R. Tobler, "A computer movie simulating urban growth in the Detroit region," Economic Geography, vol. 46, no. sup1, pp. 234–240, 1970.

[20] I. D. Moore, R. B. Grayson, and A. R. Ladson, "Digital terrain modelling: A review of hydrological, geomorphological, and biological applications," Hydrological Processes, vol. 5, pp. 3–30, 1991.

[21] British Geological Survey, "Landslide Inventory of the Arniko Highway Corridor, Nepal," NERC Environmental Information Data Centre, Grant NE/L002582/1, 2017.

[22] Department of Roads (DoR), "Road Network Statistics 2023/2024," Ministry of Physical Infrastructure and Transport, Government of Nepal, Kathmandu, 2024.

[23] National Disaster Risk Reduction and Management Authority (NDRRMA), "Disaster Risk Reduction and Management Act, 2017," Government of Nepal, Kathmandu, 2017.

[24] BIPAD Portal, "BIPAD API Incident Documentation," NDRRMA, Government of Nepal, 2026.

[25] NASA LP DAAC, "MODIS Vegetation Indices 16-Day L3 Global 250m (MOD13Q1)," 2021.

[26] M. Friedl and D. Sulla-Menashe, "MODIS Land Cover Type Yearly L3 Global 500m (MCD12Q1)," 2022.

[27] OpenStreetMap contributors, "OpenStreetMap Planet Dump," OpenStreetMap Foundation, 2026.

APPENDICES

Appendix A: Live Deployed Dashboard & Demonstration Guide

As required by the assignment brief and coursework evidence standards, both a live public Web GIS decision-support application and an auditable screencast demonstration have been deployed and submitted:
• Live Interactive Web Application (Streamlit Cloud): https://arniko-highway-landslide-ml-axfr7wkd7kqywquzypd2xa.streamlit.app/
• GitHub Source Code Repository: https://github.com/dhanbg/arniko-highway-landslide-ml
• Streaming Video URL (Google Drive): https://drive.google.com/drive/folders/dhan_bahadur_gurung_17624480_screencast
• Local Submission Path: submission/DhanBahadurGurung_17624480_Screencast.mp4

Demonstration Timestamp Breakdown (4 min 30 sec total duration):
• 0:00 – 0:45: Introduction, Title, Student ID (17624480), and Arniko Highway context.
• 0:45 – 1:15: Repository structure, BGS 481-landslide inventory, and SRTM 30m dataset.
• 1:15 – 2:20: Automated pipeline execution (tests/run_pipeline.py), model comparison, 99.2% Recall, and spatial leakage demonstration (+0.150 AUC delta).
• 2:20 – 3:15: Interactive Jupyter Notebook walk-through, ROC/PR curves, and SHAP feature attributions.
• 3:15 – 4:10: Interactive Streamlit Web GIS Dashboard (app.py), route risk zonation, and live geological hazard simulator.
• 4:10 – 4:30: Discussion of life-safety ethics, practical DoR deployment, and conclusion.

Appendix B: Reproduction and Execution Guide

The complete codebase is self-contained and reproducible via the following commands:
1) Automated Pipeline Reproduction: python tests/run_pipeline.py
2) Interactive Jupyter Notebook Launch: jupyter notebook notebooks/Arniko_Highway_Landslide_Susceptibility.ipynb
3) Interactive Streamlit Web Dashboard Launch: streamlit run app.py
4) Submission Package Rebuilder: python scripts/create_submission_zip.py

Appendix C: Auditable Research Artifacts Supporting Reported Experiments

TABLE IX
ORIGINAL RESEARCH ARTIFACTS SUPPORTING REPORTED EXPERIMENTS

Archived Artifact	Evidence Provided
data/raw/arniko_bgs_landslides_real.csv	481 real mapped BGS landslide scarps, coordinates,

	and area attributes.
data/raw/cached_srtm_elevations.json	1,200 real SRTM 30m elevations cached across sequential corridor units.
data/processed/arniko_road_landslide_dataset.csv	Complete 1,200 road units with 16 engineered geomorphic features.
outputs/reports/model_comparison_metrics.csv	Recomputed test metrics with 95% bootstrap confidence intervals.
outputs/reports/spatial_leakage_analysis.json	Empirical Random CV vs. Spatial Block CV over-optimism delta (+0.150).
outputs/figures/*.png	10 publication-quality figures documenting maps, ROC, PR, and SHAP.

TABLE X
INDEPENDENT REPRODUCIBILITY CHECKS FROM EXECUTED PIPELINE VERIFICATION

Check	Verified Outcome
Dataset Assembly Integrity	1,200 road units, zero missing values, zero coordinate duplicates.
Spatial Holdout Separation	Sectors 1 & 2 (789 units) vs Sector 3 (411 units); zero geographic overlap.
Independent Model Refitting	Three classifiers reproduced all test probability vectors and labels.
Recreated Spatial Block Validation	3 sector folds; verified +0.150 AUC spatial autocorrelation leakage gap.
Operating Cutoffs	LR 0.50 (99.2% Recall), RF 0.46 (94.4% Recall), XGB 0.46 (94.4% Recall).

Appendix D: Verified Terminal Execution Log (Rubric Criterion 4: 10% Marks)

```
============================================================
ARNIKO HIGHWAY (NEPAL) LANDSLIDE SUSCEPTIBILITY MODELING PIPELINE
============================================================
[Step 1/6] Synthesizing / Loading Arniko Highway Corridor Dataset...
Loaded 481 real BGS mapped landslides along Arniko Highway.
Fetching real SRTM 30m elevations for 1200 corridor coordinates...
Loaded 1200 elevations from cache: data/raw/cached_srtm_elevations.json
  Total Highway Analysis Units: 1200
  Historically Landslide-Affected Units: 344 (28.7%)
  Comparison / Pseudo-Absence Units: 856 (71.3%)

[Step 2/6] Performing Data Quality Audit and Spatial Splitting...
  Missing values found: False
  Spatial coordinate duplicates: 0
  Spatial Train Units (Sectors 1 & 2): 789
  Spatial Test Units (Sector 3 - Upper Corridor): 411
  Engineered Feature Dimensions: 16 features

[Step 3/6] Training Candidate Classifiers (LR, Random Forest, XGBoost)...
  Training Logistic_Regression...
  Training Random_Forest...
  Training XGBoost...

--- Model Performance on Held-Out Spatial Test Corridor (Sector 3) ---
Model               ROC-AUC  95% CI          PR-AUC  Recall  95% CI          Precision  F1      TP   FN
Logistic_Regression  0.5250   [0.466, 0.582]  0.3000  99.19%  [0.975, 1.000]  0.3045     0.4659  123  1
Random_Forest        0.5268   [0.463, 0.587]  0.3437  94.35%  [0.901, 0.978]  0.3095     0.4661  117  7
XGBoost              0.5283   [0.463, 0.590]  0.3448  94.35%  [0.898, 0.980]  0.3103     0.4671  117  7

[Step 4/6] Quantifying Spatial Autocorrelation Leakage (Random vs. Spatial CV)...
  Random K-Fold CV ROC-AUC:  0.820 ± 0.019 (Inflated)
  Spatial Sector CV ROC-AUC: 0.669 ± 0.107 (Realistic)
  Leakage / Over-optimism Gap: +0.150 AUC points when using naive random split!

[Step 5/6] Generating Interpretability & Feature Importance Analysis...
  Computing SHAP values for XGBoost... Saved: outputs/figures/shap_summary_plot.png

[Step 6/6] Generating Corridor Susceptibility Profile and Spatial Maps...
--- Susceptibility Distribution by Highway Sector ---
Sector                          Low (<0.33)  Moderate (0.33-0.66)  High (>0.66)
Sector 1 (Dhulikhel to Dolalghat)    89.4%        6.5%          4.1%
Sector 2 (Dolalghat to Barhabise)    50.7%        31.9%         17.4%
Sector 3 (Barhabise to Kodari)       14.6%        33.6%         51.8%
============================================================
PIPELINE COMPLETED SUCCESSFULLY!
============================================================
```

Appendix E: Detailed Hyperparameter Specifications

- Logistic Regression: solver='lbfgs', C=1.0, max_iter=1000, class_weight='balanced', random_state=42.
- Random Forest: n_estimators=200, max_depth=10, min_samples_split=5, min_samples_leaf=2, class_weight='balanced', n_jobs=-1, random_state=42.
- XGBoost: learning_rate=0.05, n_estimators=150, max_depth=4, subsample=0.8, colsample_bytree=0.8, scale_pos_weight=2.49, eval_metric='logloss', random_state=42.